



# Convolutional Neural Networks

---

Machine Learning Decal

Hosted by Machine Learning at Berkeley

# Agenda

Motivation

Convolutions

Convolutions in Neural Networks

Demos

Questions

## Motivation

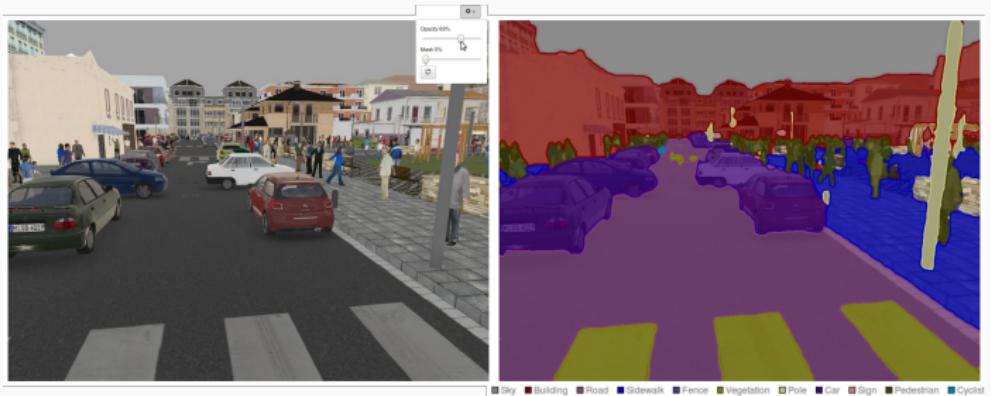
---

# Applications of CNNs



- Computer vision
  - Face recognition
  - Scene labeling
  - Image segmentation
  - Image classification
  - Action recognition
  - Human pose estimation
  - Playing Games (Atari games, Pacman)
  - Medical image analysis

# Image Segmentation



# Image Generation



# Image Style Transfer

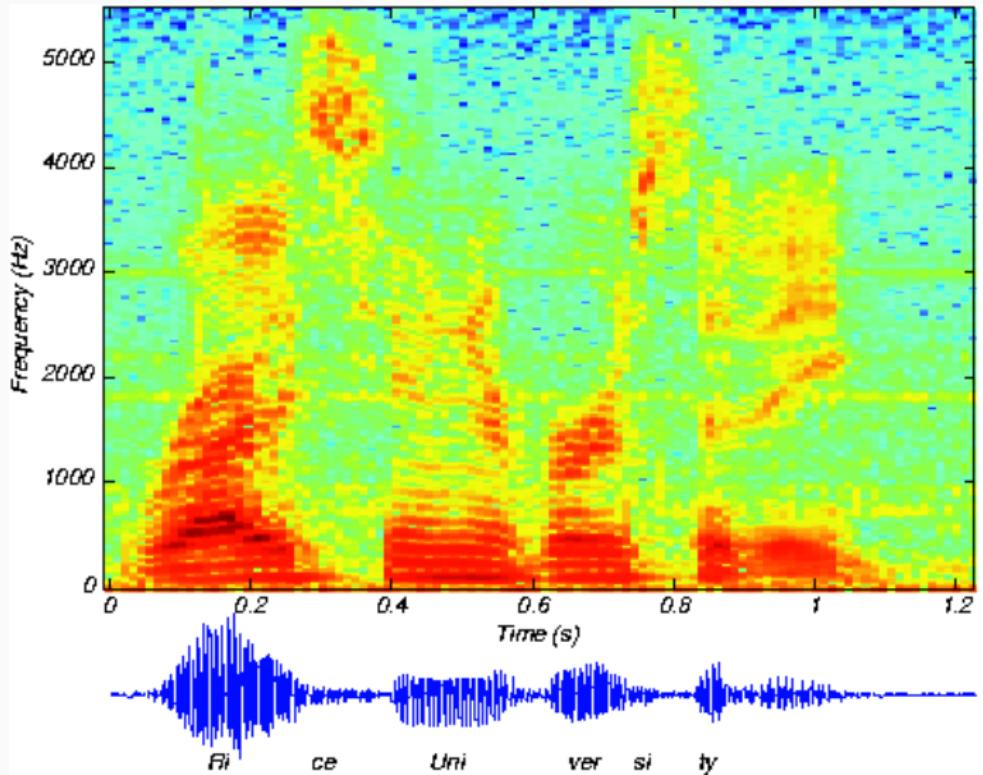


# Another Style Transfer Example



Cycle GAN: <https://github.com/junyanz/CycleGAN>

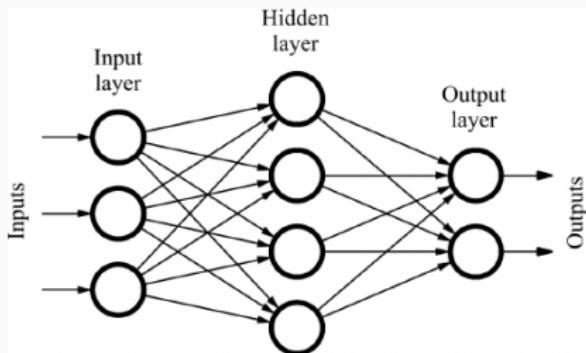
# Audio Processing



# Industry use of CNNs?



# What's wrong with feed-forward NNs?



The curse of dimensionality.

# Problem: high-dimensional image data

- If we want a neural network to take as input  $256 \times 256$  images, the input layer will have 65,536 neurons
- ... if we have one hidden layers with 1,000 neurons, we're already at 65 million parameters
- ... if we have two hidden layers with 15,000 neurons each, we're at over a **billion** parameters

Basic Idea: Share parameters for different parts of input.

# Convolutions

---

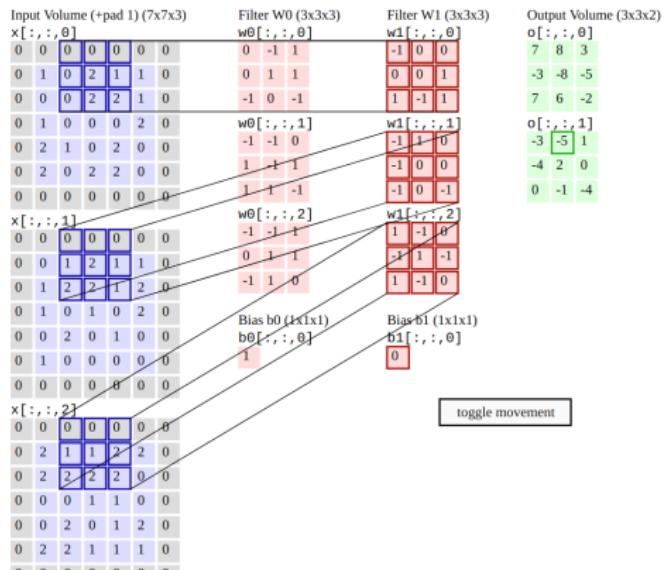
Intuition: sliding dot product.

Suppose  $f$  is a length  $N$  vector.

$$(g * f)[n] = \sum_{i=0}^N g[n + i]f[i]$$

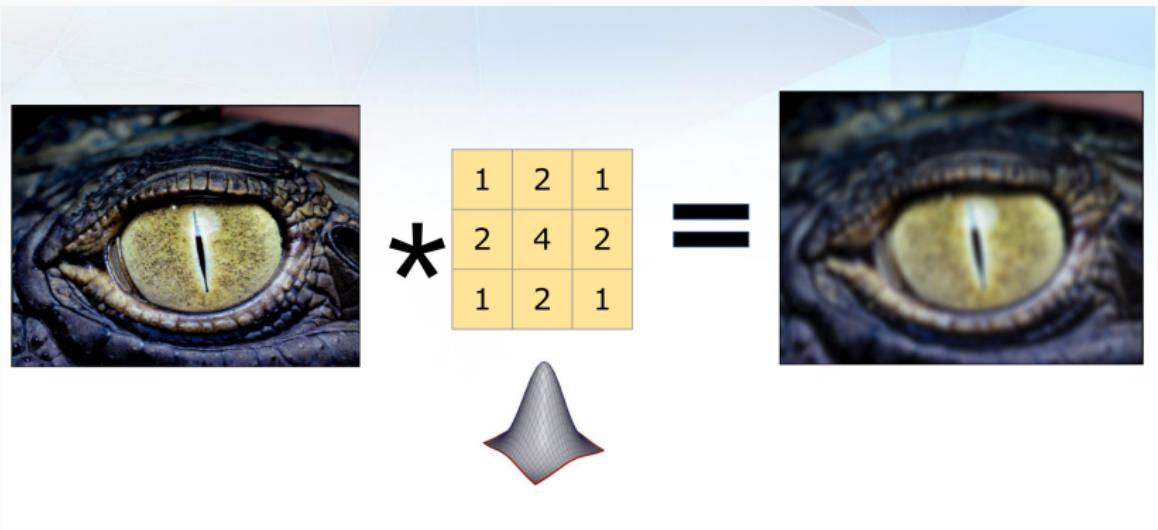
We can think of  $f$  as a filter (more on this later).

# 2D Convolution: Most Common in CNNs



<http://cs231n.github.io/convolutional-networks/>

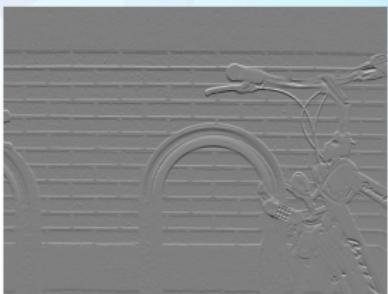
# Convolutions as filters



# Convolutions as filters cont.

 $\ast$ 

1	2	1
0	0	0
-1	-2	-1

 $=$ 

# Demo

This will account for exactly  
0% of your grade in this class.

If  $f = [5, 1, 2]$  and  $g = [3, 1, 4, 5, 6, 0, 0, 1]$ , What is  
the 4 element of the output vector of  $f * g$ ?

# Answer

$$f * g[4] = 5 * 5 + 1 * 6 + 0 * 2$$

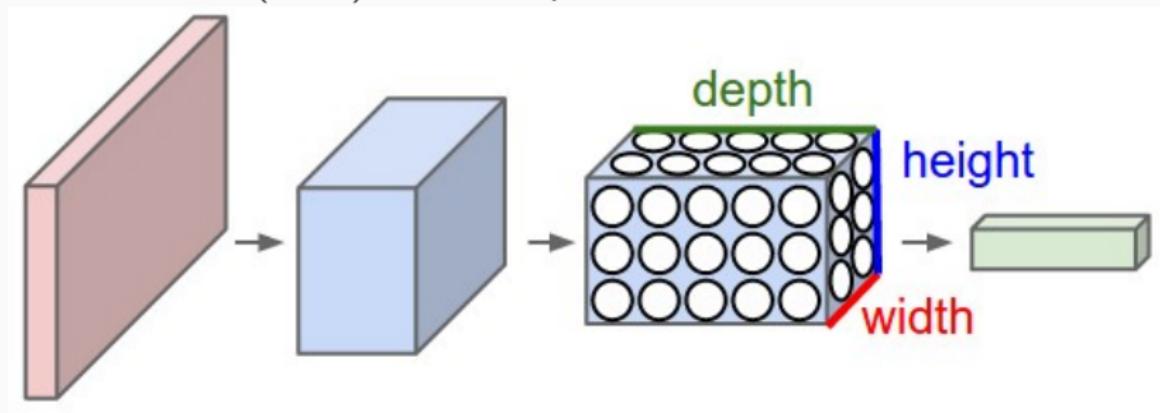
# **Convolutions in Neural Networks**

---

CNN = Learned convolutional filters.

# What does our input look like?

Let's say we have a 32x32 color image. There are then 3 channels for the colors (RGB). So our input is a 3D volume of size 32x32x3



# What do our filters look like?

3D filters !!!

Our filters will then be  $n \times n \times 3$ , where  $n$  is a hyperparameter.

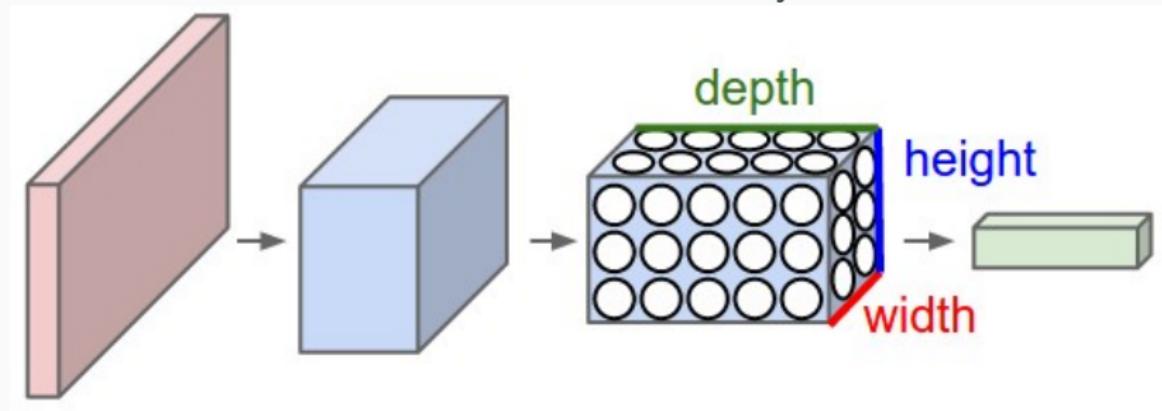
We can also have multiple filters per layer.

Input Volume (+pad 1) (7x7x3)	Filter W0 (3x3x3)	Filter W1 (3x3x3)	Output Volume (3x3x2)
$x[:, :, 0]$	$w0[:, :, 0]$	$w1[:, :, 0]$	$o[:, :, 0]$
0 0 0 0 0 0 0	0 -1 1	-1 0 0	7 8 3
0 1 0 2 1 1 0	0 1 1	0 0 1	-3 -8 -5
0 0 0 2 2 1 0	-1 0 -1	1 -1 1	7 6 -2
0 1 0 0 0 2 0	w0[:, :, 1]	w1[:, :, 1]	$o[:, :, 1]$
0 2 1 0 2 0 0	-1 -1 0	-1 1 -1	-3 -5 1
0 2 0 2 2 0 0	1 -1 1	-1 0 0	-4 2 0
0 0 0 0 0 0 0	1 -1 -1	-1 0 -1	0 -1 -4
$x[:, :, 1]$	$w0[:, :, 2]$	$w1[:, :, 2]$	
0 0 0 0 0 0 0	-1 -1 1	1 -1 0	
0 0 1 2 1 1 0	0 1 1	-1 1 -1	
0 1 2 2 1 1 2	-1 1 0	1 -1 0	
0 1 0 1 0 2 0	Bias b0 (1x1x1)	Bias b1 (1x1x1)	
0 0 2 0 1 0 0	$b0[:, :, 0]$	$b1[:, :, 0]$	
0 1 0 0 0 0 0	1	0	
0 0 0 0 0 0 0			
$x[:, :, 2]$			
0 0 0 0 0 0 0			
0 2 1 1 2 2 0			
0 2 2 2 2 0 0			
0 0 0 1 1 0 0			
0 0 2 0 1 2 0			
0 2 2 1 1 1 0			
0 0 0 0 0 0 0			

toggle movement

# What will the output of a layer look like?

Our output will also be 3D of shape  $(32 \times 32 \times m)$  where  $m$  is the number of filters in the layer.



# How do we update filters?

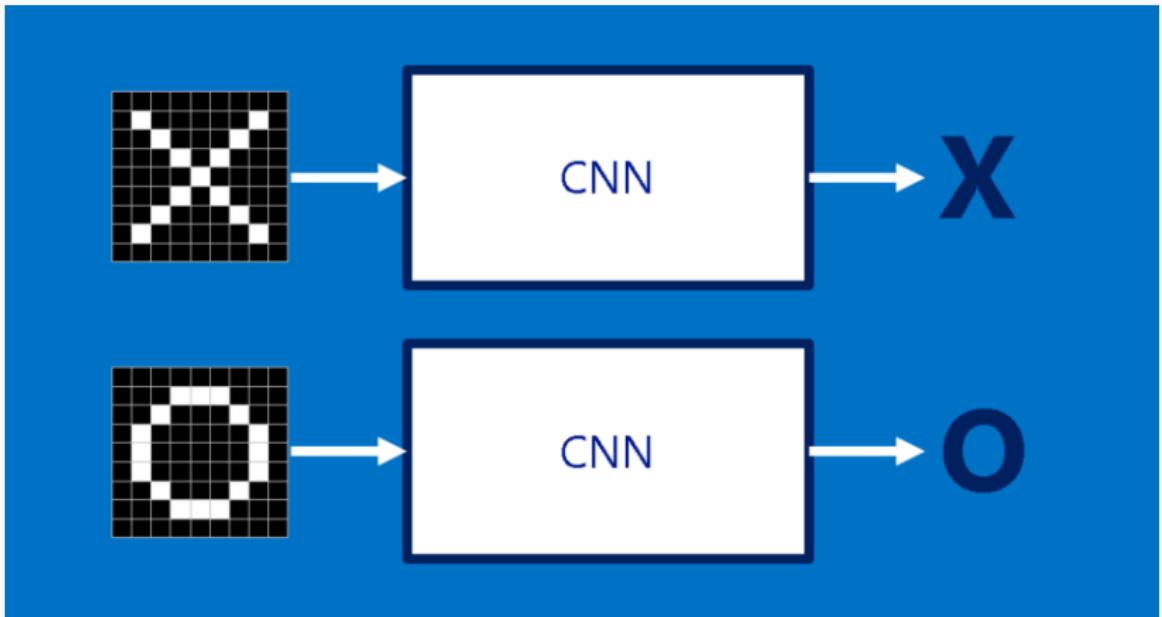
$X_{11}$	$X_{12}$	$X_{13}$
$X_{21}$	$X_{22}$	$X_{23}$
$X_{31}$	$X_{32}$	$X_{33}$

$X_{11}\partial h_{11}+$ $X_{12}\partial h_{12}+$ $X_{21}\partial h_{21}+$ $X_{22}\partial h_{22}$	$X_{12}\partial h_{11}+$ $X_{13}\partial h_{12}+$ $X_{22}\partial h_{21}+$ $X_{23}\partial h_{22}$
$X_{21}\partial h_{11}+$ $X_{22}\partial h_{12}+$ $X_{31}\partial h_{21}+$ $X_{32}\partial h_{22}$	$X_{22}\partial h_{11}+$ $X_{23}\partial h_{12}+$ $X_{32}\partial h_{21}+$ $X_{33}\partial h_{22}$

$\partial h_{11}$	$\partial h_{12}$
$\partial h_{21}$	$\partial h_{22}$

Derivative Computation (Backward pass) since pictures speak more than words

# Let's look at a toy example!



- How would a CNN classify X's vs O's?

# Would you classify this as an X?

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	
-1	-1	1	-1	-1	-1	1	-1	-1	
-1	-1	-1	1	-1	1	-1	-1	-1	
-1	-1	-1	-1	1	-1	-1	-1	-1	
-1	-1	-1	1	-1	1	-1	-1	-1	
-1	-1	-1	1	-1	1	-1	-1	-1	
-1	-1	1	-1	-1	-1	1	-1	-1	
-1	1	-1	-1	-1	-1	-1	1	-1	
-1	-1	-1	-1	-1	-1	-1	-1	-1	

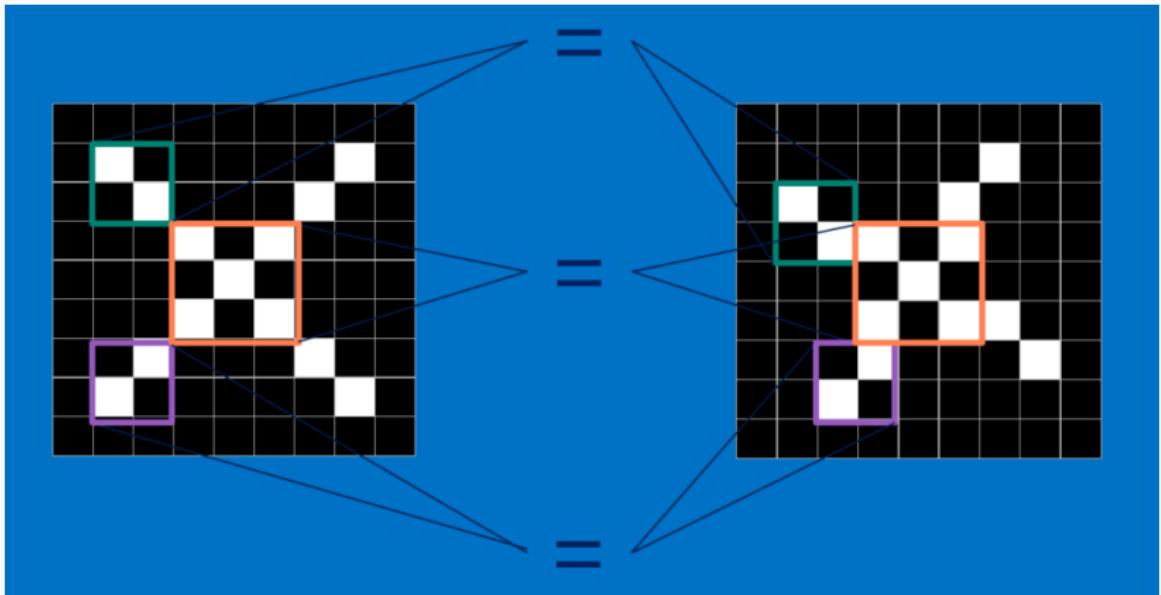
?



-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

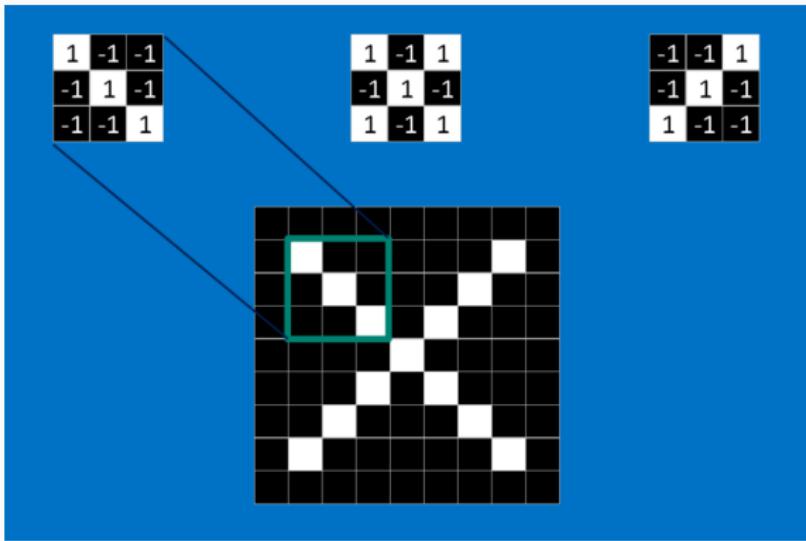
- A computer can't easily distinguish this example as an X.. so what can we do?

# What is a feature?



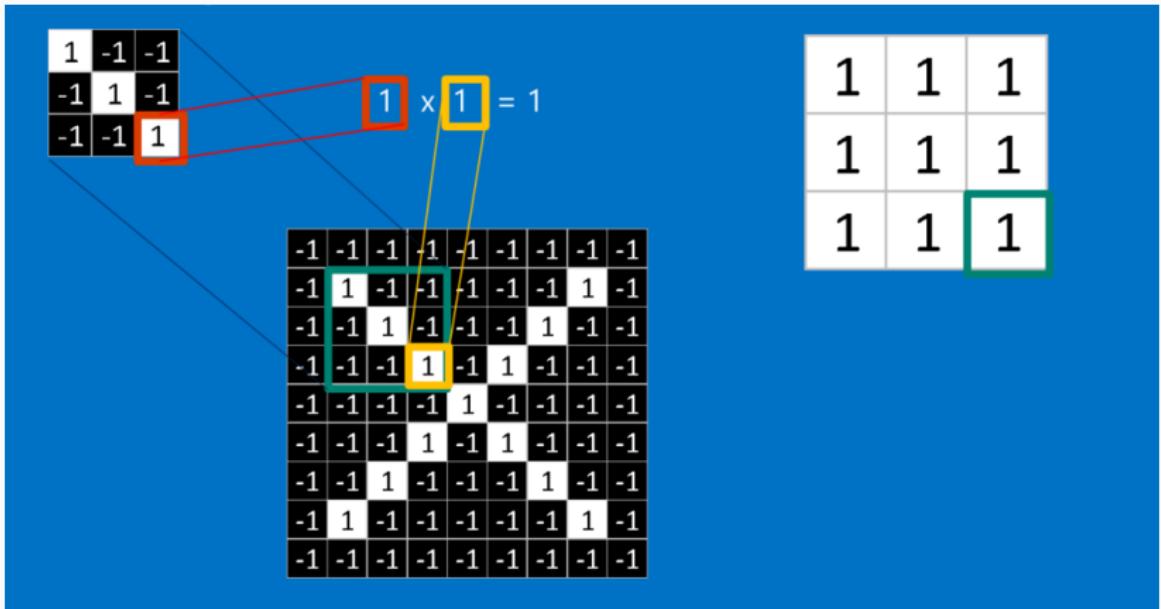
- We can break down the image into smaller images and look for patterns!
- Can we see here the similar patterns?
- These diagonals and crosses can be thought of as features!

# Convolutional Filters



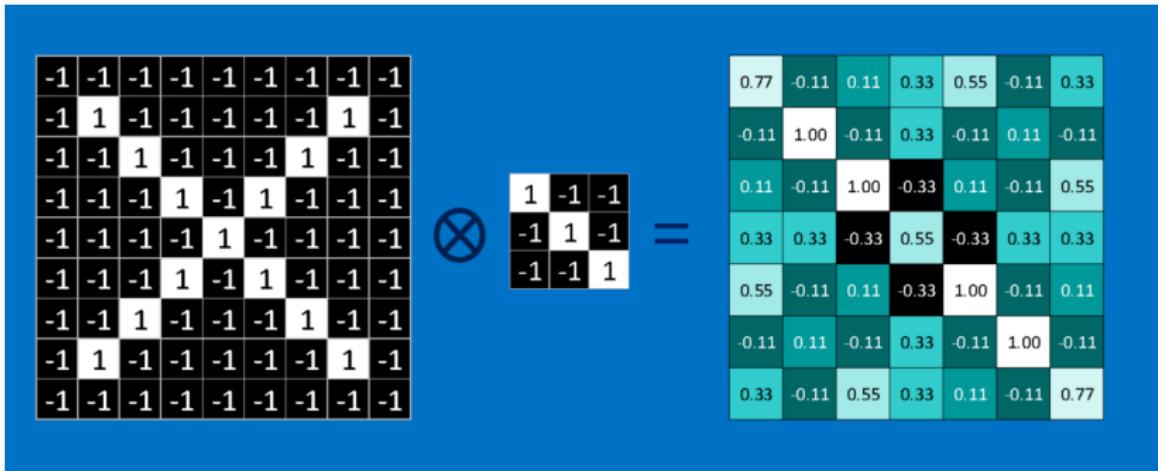
- Let's break this up into smaller problems... what about smaller patterns!

# CNN Layers: Convolutions



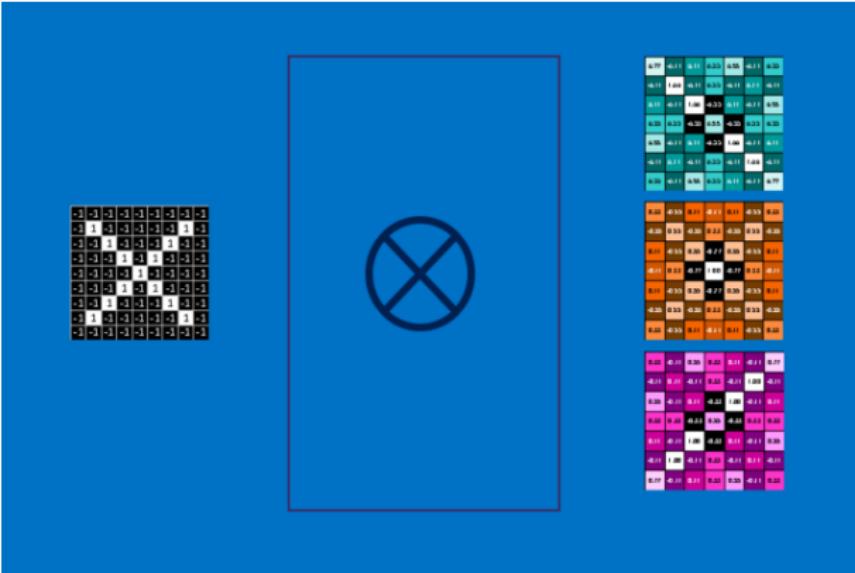
- Now let's compare our smaller pattern with the image.
- This is the basic idea behind convolutions!

# CNN Layers: Convolutions



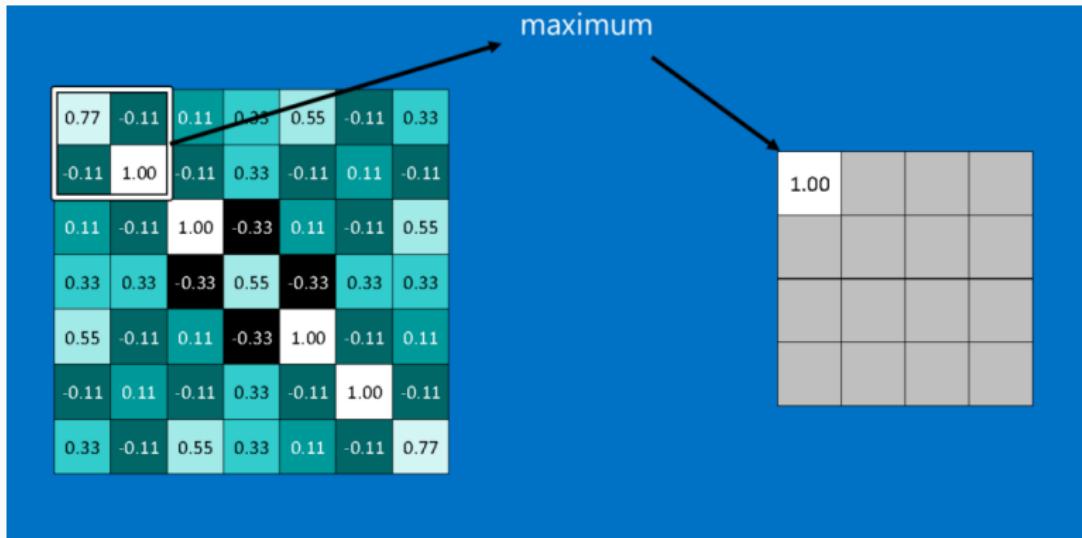
- Now let's apply this process to every part of the board.
- We can look at the new board and call it a filtered version of this specific feature
- Shows us where on the image the feature matches the most

# CNN Layers: Convolutions



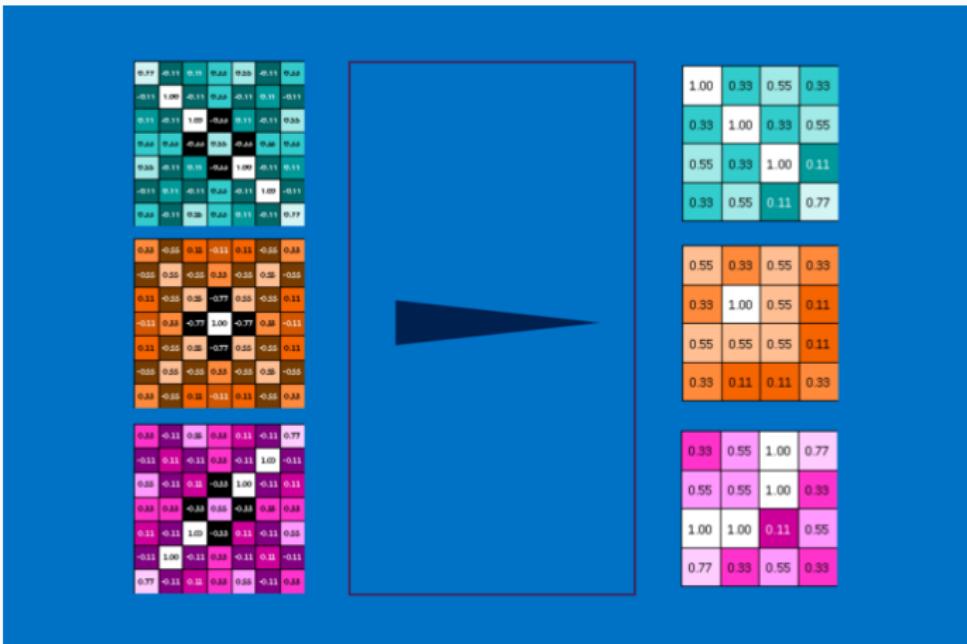
- Now lets apply this process to every part of the image for different features.

# CNN Layers: Pooling



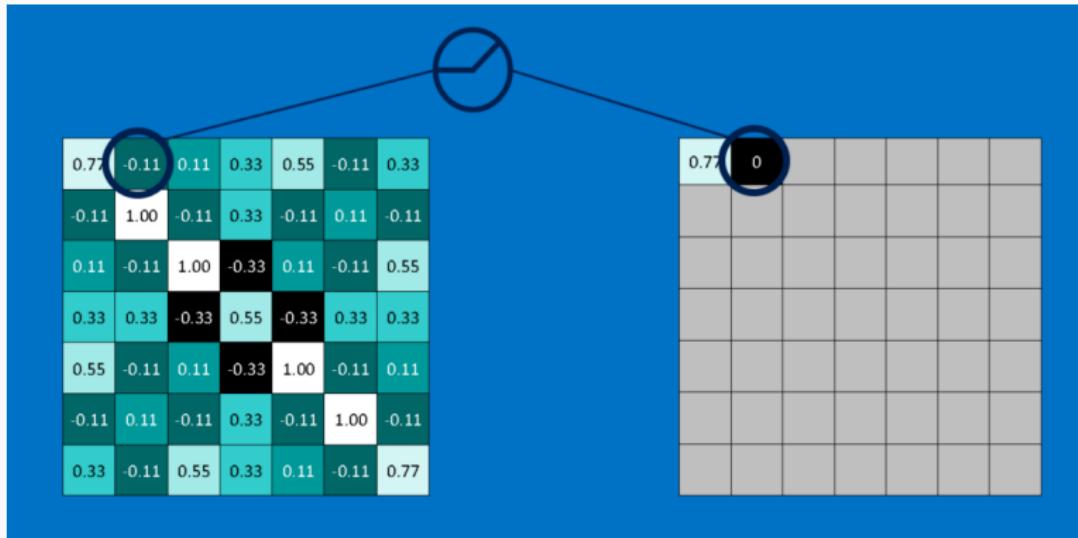
- If CNN's are really deep, doing that math for every possible pixel gets to be quite computationally expensive
- That's why we have pooling layers, that down sample our image!

# CNN Architecture



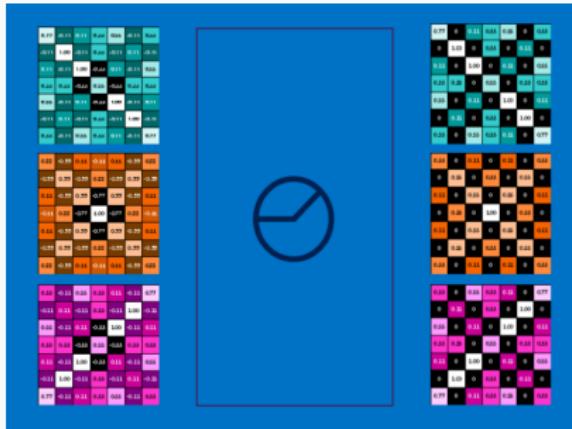
- By applying these pooling layers, we try to capture the patterns and get rid of extra unneeded information.

# CNN Layers: Activation



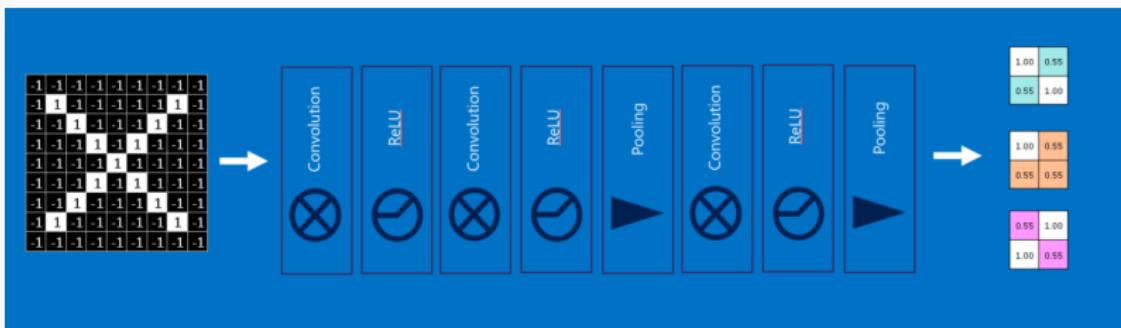
- Remember those activation functions you learned from the last lecture? We use them here too!
- Applying RELU we simply get rid of all the negatives!

# CNN Pooling Layers

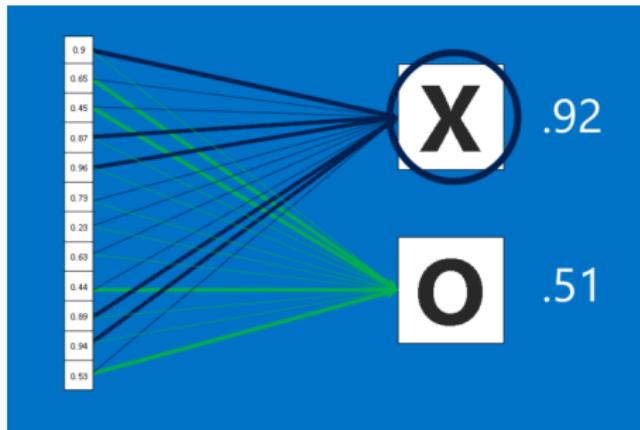


- The idea behind this, is to help keep any one number stuck at 0, or block up towards infinity
- This may not seem like a big issue in our toy example, but in much larger more complicated CNN's having stable values is important!

# Let's put it all together!



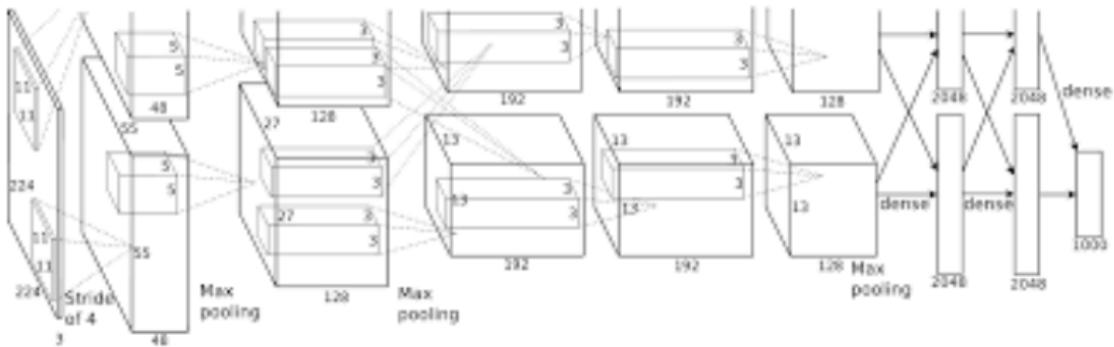
# CNN: Fully Connected Layer



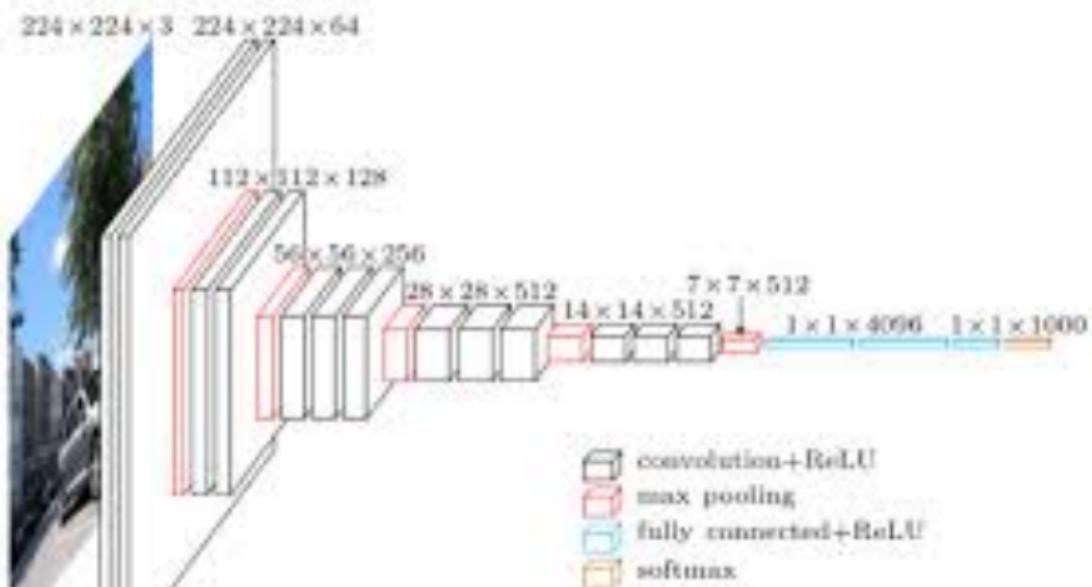
- At the end we use a fully connected layer to hold a vote if the image is either an X or O
- Each value has a weight though, so not all values contribute equally
- This goes back to the ideas you learned last week!

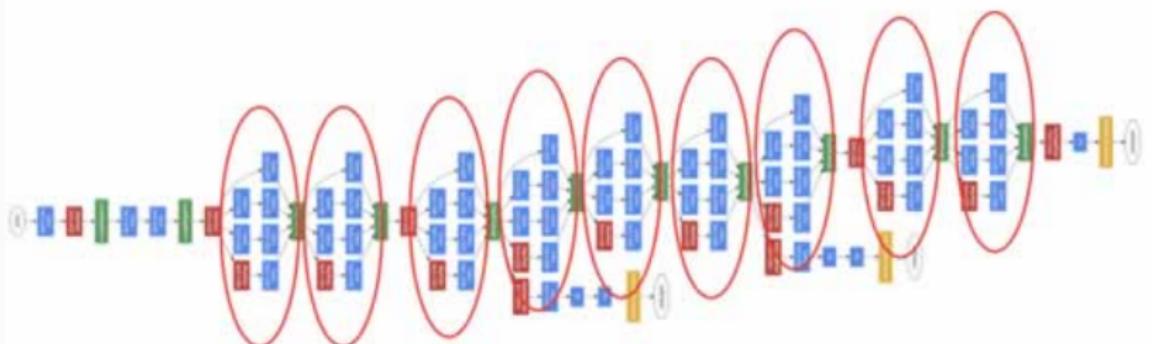
- "The 9 Deep Learning Papers You Need to Know About" (<https://adeshpande3.github.io/adeshpande3.github.io/The-9-Deep-Learning-Papers-You-Need-To-Know-About.html>)
- Many of these models are used as benchmarks in deep learning
- There exist many pre-trained versions of the above models that can be leveraged for a new task (this is called **transfer learning**)
- Let's look at a few...

# AlexNet (2012)



# VGG-16 (2014)



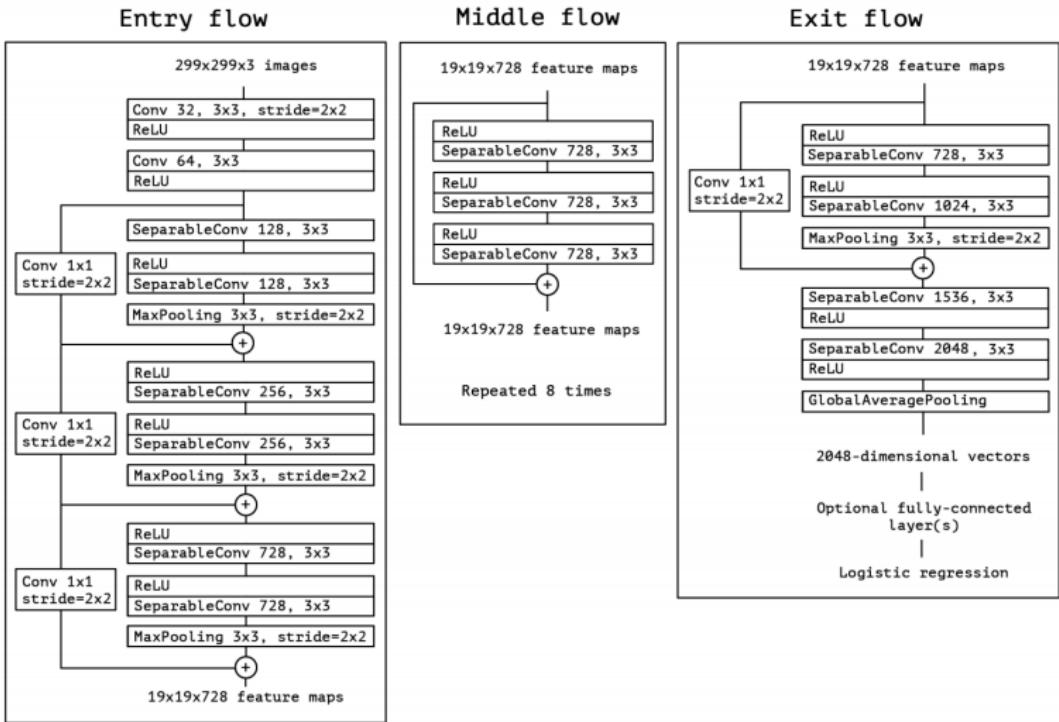


## 9 Inception modules

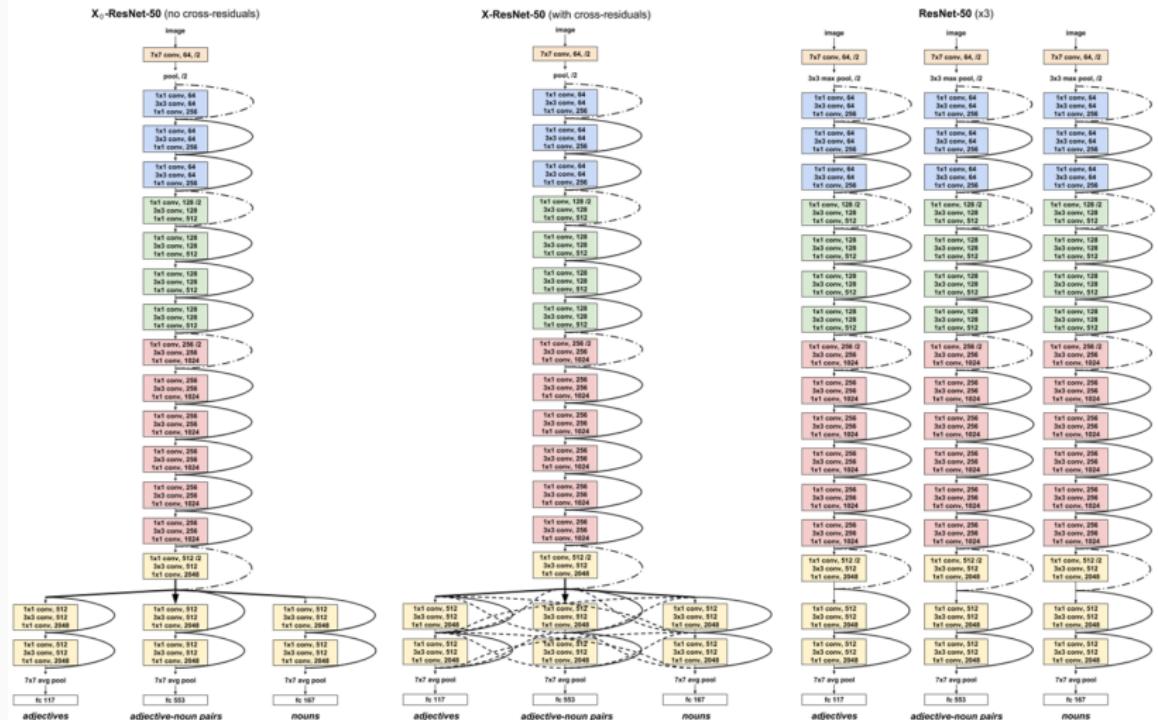
Network in a network in a network...

Convolution
Pooling
Softmax
Other

# Xception (2016)

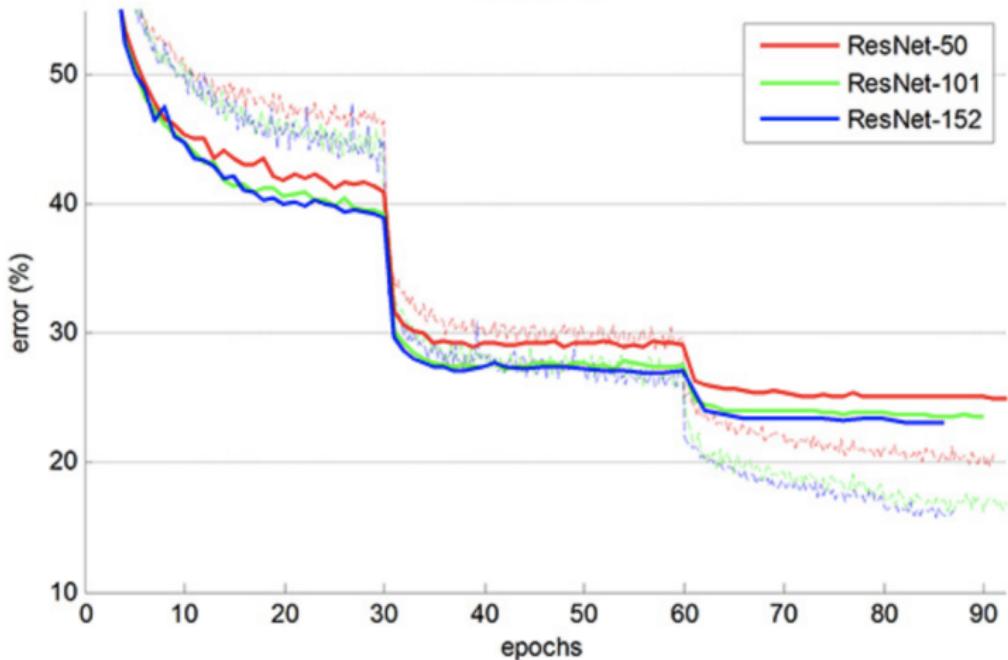


# ResNet-50 (2015)



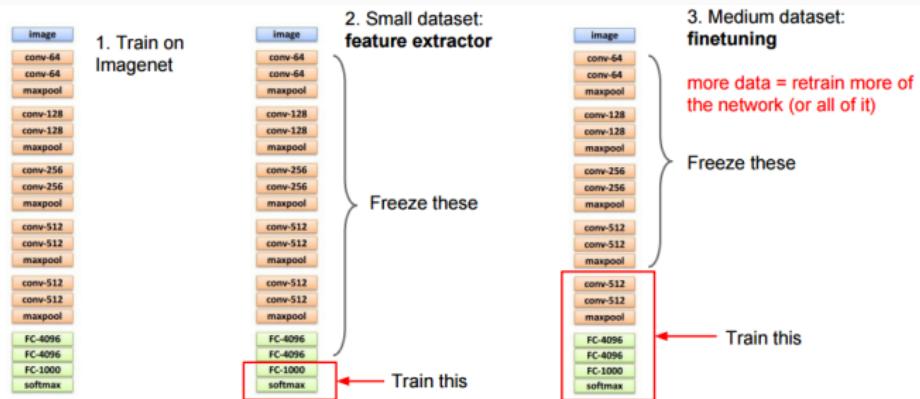
# ResNet-152 (2015)

ImageNet-1k



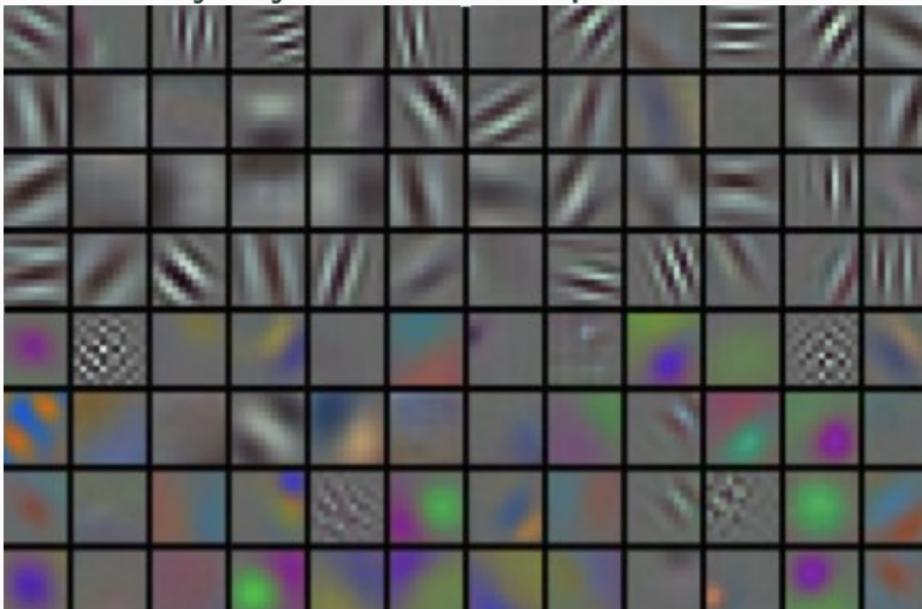
- In practice, we don't train an entire ConvNet from scratch
- One can pretrain the CNN on a large dataset like ImageNet and then finetune the ConvNet so that it does well on our dataset.
- There are several scenarios when we might want to use transfer learning.
- Let's look at a few...

# How Can We Transfer Learn?



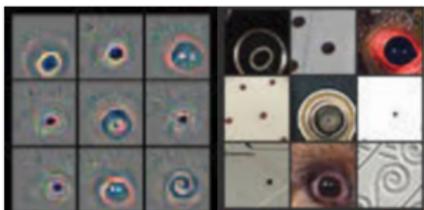
# Why does it work?

Early layers learn simple features

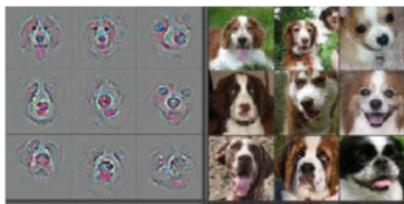


# Why does it work?

Later layers learn more abstract features



Layer 2 of AlexNet



Layer 4 of AlexNet

## Demos

---

# CNN demo: FashionMNIST

- Training networks using GPUs
- Network visualization and interpretation
- Other types of neural networks: auto-encoders, Neural Turing machines, Generative adversarial networks, ...

## Questions

---

# Questions?